

# README

## SECTION 1: FILES

---

This readme contains information about how to use the files downloaded in this folder. The folder should contain:

- `README.pdf` – This file, explains contents of export and how to use the files.
- `Dockerfile` – Dockerfile, can be used to setup an environment similar to that of the application.
- `requirements.txt` - Python requirements file, used by the Dockerfile, could be used in another environment setup methodology.
- `start_notebook.sh` - Script to startup a jupyter notebook server. **Not** meant to be run outside of docker container.
- `UsageExamples.ipynb` - Jupyter Notebook which demonstrates usage of the files & model (see section 4).
- `project_#_tfidf_matrix.pkl` – TFIDF (see section 2) matrix of the uploaded data.
- `project_#_training_#.pkl` – Model (see section 3), trained on the most recent labeled data
- `project_#_labeled_data.csv` – All labeled data, with the original text, unique ID, and assigned label.
- `project_#_labels.csv` – Mapping between label name and ID.
- `project_#_vectorizer.pkl` – Fitted TFIDF Vectorizer (see section 2). This can be used to fit new data to the same format as the training data (see section 4 for examples)

**NOTE:** All models and TFIDF files use Scikit-Learn version 0.19.0. If you are using a different version, be careful! It is recommended to use the provided requirements.txt to prepare an environment or use the supplied Dockerfile to create one for you.

## SECTION 2: TFIDF

---

### A. What is TFIDF?

Term Frequency Inverse Document Frequency is a way of quantifying text data. For each document (piece of text, in this case one tweet, news article, etc.) a count is created for the number of times each term (word) appears. This count is then scaled down by the amount of times the word appears in all of the documents (so common words like “the” are given a lower score). The result is a set of numerical features representing each document by the words in it.

Note that this method of quantifying text is a bag of words approach, and does not take the location or co-occurrence of words or phrases within a document into account.

## B. File format

The text data loaded into SMART is preprocessed with Scikit-Learn's [TFIDFVectorizer](#) with the parameters:

- `max_df`: 0.995 (only keep those terms with document frequency lower than this value)
- `min_df`: 0.005 (only keep those terms with document frequency higher than this value)
- `stop_words`: English (Automatically remove words like “the”, “at”, “and”, etc.)

The data is then placed into a dictionary of unique id to values (ex: { “tweet1” : [0, 0.5, 0.8, .... ], “tweet2”: [0.56, 0.0, 0.99, .....] } ). It is saved as a pickle file with the .pkl ending.

# SECTION 3: THE MODEL

---

## A. Source

The models used in this application are build using Scikit-Learn libraries. The four options are:

- [Logistic Regression](#)
  - Parameters: *classweight: balanced, solver: lbfgs, multiclass: multinomial*
- [Support Vector Machine](#)
  - Parameters: default
- [Random Forest](#)
  - Parameters: default
- [Gaussian Naïve Bayes](#)
  - Parameters: default

## B. File Format

All models are saved as pickle (.pkl) files through Scikit-Learn's joblib library.

# SECTION 4: HOW TO RUN

---

First you should set up a python3 environment with the libraries (and versions) specified in requirements.txt. You can use whatever method you prefer, however if you have Docker setup and would prefer to utilize the provided Dockerfile and jupyter notebook examples. After navigating to the directory containing all of these files, follow the instructions below:

```
$ docker build -t smart_export -f Dockerfile .
$ docker run -it -p 8888:8888 -v $(pwd):/code/ smart_export ./start_notebook.sh
```

**NOTE:** If you are not in a unix-like environment or do not have the `pwd` (print working directory) command for whatever reason, then replace `$(pwd)` with the **absolute** path to the current directory.

Then go to [localhost:8888](http://localhost:8888) and click on the `UsageExamples.ipynb` to launch the notebook which demonstrates usage. You can create new notebooks for further analysis and they will be saved in this directory.

Subsections A and B walk through the code in a text-format. If you are using the UsageExamples notebook then you do not need to read further.

## A. Preprocessing new data for the model to predict

The following sample code can be used to preprocess new data and get it in a format that is usable by the saved model.

```
import pickle

with open(<<project_#_vectorizer.pkl>>,"rb") as tfidf_vectorizer:
    # load the vectorizer
    tfidf_transformer = pickle.load(tfidf_vectorizer)

new_data = ["This is new data","that needs to be formatted", "In the same way that the other data was before"]

# apply it to new data
transformed_data = tfidf_transformer.transform(new_data)
```

**Note:** to see what words the fields in the transformed TFIDF matrix are based off of, just call `print(tfidf_transformer.vocabulary_)`. This will print a dictionary of the words that were used and their index in the transformed tfidf\_matrix.

## B. Applying the model to predict new data

The following sample code can be used to read in the zipped files and get predictions for the unlabeled data. (The `<<>>>` and `#` are placeholders for your file names and project number).

**NOTE:** the model will predict labels as a number. The project\_#\_labels.csv gives the mapping between label text and ID.

```
import pandas as pd
import pickle
from sklearn.externals import joblib

# read in the TFIDF matrix and the labeled data
labeled_frame = pd.read_csv(<<project_#_labeled_data.csv>>)
with open(<<project_#_tfidf_matrix.pkl>>,"rb") as tfidf_file:
    tfidf_dict = pickle.load(tfidf_file)

# Subset the TFIDF matrix by the unlabeled data and get as 2D list
labeled_ids = labeled_frame["ID"].tolist()
unlabeled = [tfidf_dict[key] for key in tfidf_dict if key not in labeled_ids]

# read in the model from the pickle file
model = joblib.load(<<project_#_training.pkl>>)

# apply the model to the unlabeled data
predictions = model.predict(unlabeled)

# print the results
print(predictions)

# open the label dictionary to translate label ID's to their corresponding text
labels_frame = pd.read_csv(<<project_#_labels.csv>>)
label_dict = labels_frame.set_index("Label_ID").to_dict()["Name"]

# get the predictions as actual labels
predictions = [label_dict[pred] for pred in predictions]
print(predictions)
```